

Introduction to Algorithms




Algorithm Analysis

Learning Objectives

- **Overall:** To promote individual learning and **analytical thinking skills**.
 - More detailed objectives:
 1. Be able to **prove correctness** of algorithms.
 - Some people spend their whole career to prove something.
 2. Have deep understanding of differences and advantages of **iterative and recursive algorithms**.
 3. Have a knowledge base of (many) **existing algorithms**.
 4. Understand **speed vs. space trade-off**, importance of **data-structures, and data preprocessing**.
 - Search for an item: hash, tree, and traverse all.
 5. Be able to **design new algorithms**.
 - Correctness.
 - Efficiency.
-
-

Overview

- Acknowledgement: Dr. Zhang Wei, Dr. LIU Fang, Dr. Donny - course materials
- Appointment by email:
 - Pohkok.Loo@singaporetech.edu.sg
 - AF_Alfred.Whang@singaporetech.edu.sg
- References:
 - **Algorithm Design**, Jon Kleinberg and Eva Tardos.
 - **Introduction to Algorithms**, Thomas Cormen, Charles Lieserson, *et al.*
 - Introduction to the Design and Analysis of Algorithms, Anany Levitin.
 - Course materials and activities



Source: <https://kwize.com/quote/9797>; <https://eden.uktv.co.uk/animals/birds/article/weaver-birds/>

Tentative Dates and Grading Policy

- **Marks:**
 - 28% - Assignment (group of 5/6 students)
 - 12% - Tutorial participation/quiz - no makeup
 - 30% - Midterm
 - 30% - Final Exam
 - **Part 1: week 1 to 6 (Loo)**
 - 1 assignment (group 14%), week 5, Due on 25/9/23 23:59
 - 5 tutorial participation/quiz (6%), during each tutorial
 - 1 midterm (30%), week 6, 3/10/23 2-3pm
 - **Part 2: week 8 to 13 (Alfred)**
 - 1 assignment (group 14%), week 12, Due on 13/11/23
 - 5 tutorial participation/quiz (6%), during each tutorial
 - 1 Final (30%), week 13, 21/11/23 (TBA)
-
-

Tentative Dates and Grading Policy

- Blended learning
 - All lecture content/video/zoom recording will be posted before the tutorial.
 - The key face to face session is during the tutorial where you will work on tutorial questions, ask questions and complete quizzes.
 - Enrol yourself in xSit follows the below guide. Minor variation may occur due to dynamic enrolment add/drop situation.
 - P1: 5 groups of 6 students, 2 groups of 5 students
 - P2: 4 groups of 6 students, 3 groups of 5 students
 - P3: 5 groups of 6 students, 2 groups of 5 students
 - P4: 6 groups of 6 students, 1 groups of 5 students
 - Can only join students within the same tutorial group
-
-

Tentative Dates and Grading Policy

- Peer review
 - If a group with 5 students received 72 for the assignment, then

Student	Contribution	Final score
S1	80/100	$80/120 \times 72 = 48.0$
S2	95/100	$95/120 \times 72 = 50.5$
S3	95/100	$95/120 \times 72 = 50.5$
S4	110/100	$110/120 \times 72 = 66.0$
S5	120/100	$120/120 \times 72 = 72.0$
	Max=120	

- Be careful when allocating percentage for each member. It should not vary too much.
- A student contribution = average contribution given by all other members and yourself.

Stable Matching: Example

- History: A self-enforcing college admission process, or job recruiting process.
 - Each student has his/her preference list of the universities.
 - Each university has her preference list of the students.
 - Normal that the preferences do not align perfectly.
 - How to make both sides satisfied?



Stable Matching: Example

- **Example** to be more intuitive (in time order):
 - (May 1) Jon accepted the offer from Apple.
 - (May 2) Google offers Jon and Jon likes Google more.
 - (May 3) Jon retracts his acceptance of Apple and accepts Google's offer. (Apple has one free slot now.)
 - (May 4) Apple soon offers David and he likes Apple more.
 - (May 5) David retracts his Facebook offer and accepts Apple's offer. (Facebook has one free slot now.)
 - (May 6) Facebook moves to X, and X turns down Y and accepts Facebook.
 - (May 7) Y moves to Z ... Situation gets **out of control**.

Stable Matching: Example

- Things can be different, possibly even worse.
 - Jon, David and the others have communication channels.
 - They hear something and accordingly do something.
 - The **network** of the communication will be **fully connected** with around $O(n^2)$ links.
 - Suppose 50 students, 50 companies, then we have 10,000 communication links.
 - Each link can trigger a change of the offer distribution.
 - **Chaos happens**. Everyone tired and unhappy.
- Basic issue?
 - People cannot act in their **self-interest**, otherwise **system may crash**.

Stable Matching: Example

- Basic issue?
 - People cannot act in their **self-interest**, otherwise **system may crash**.
- Our desired process – **stable situation**:
 - Use **intelligence to guide** the process.
 - **Idea: self-interest prevents retracting offers**.
 - **Company**, with its offers all being accepted, gets a call from another applicant
 - -> “No, we prefer each applicant we have accepted than you, so we **will not change**.” **Chaos avoided**.
 - **Student**, with an offer, gets a call from a university.
 - -> “No. **I am happy** where I am.” No chaos.

Stable Matching


- Question: given employer E and applicant $A \notin E$, can we have:
 - E prefers her accepted applicants than A (another applicant).
 - A prefers his/her current employer over working for E .
- Obj: design algorithm -> produce **stable situations**.
- Similar application: residents and hospitals matching.
 - Even a decade earlier than the employee and employer one.
 - Still widely used today in USA.
- Intuition: personal interest & system interest.
- Problem formulation -> algorithm design -> algorithm analysis.

Problem Formulation

- Problem: As **clean** as possible. As **simple** as possible. Improve **step by step**.
 - Apply to programming as well.
- Complex situation: a **company can hire multiple employees**.
 - 1-to-many relationship.
- Simple: **marriage problem**.
 - A man can only marry a woman.
 - A woman can only marry a man.
 - **1-to-1 relationship**.
- **Math definition**.
 - n men: $M = \{m_1, m_2, \dots, m_n\}$.
 - n women: $W = \{w_1, w_2, \dots, w_n\}$.
 - **Matching**: S , a set of pairs, pair $(m, w) \in S$ has a man $m \in M$ and a woman $w \in W$.
 - Requirement: each man/woman only **appears at most once in S** .
 - **Perfect matching**: everyone married & married to one person. Or $|M| = |W| = |S| = n$.


Problem Formulation

- **Preference**: each man m ranks all women. The order is m 's preference list.
 - m prefers w to w' if he ranks w higher than w' .
 - Ties not allowed.
 - Each woman w also ranks all the men, also her preference list.
- Perfect matching all we want? (All men and women in the matching)
- **Unstable pair** (m, w) and (m', w') in a perfect matching S :
 - m prefers w' .
 - w' prefers m .
 - (m, w') is called an **instability**. The pair is not in S .
 - m and w' prefer each other but they are assigned to a different person
 - Objective: find a **perfect matching, without any instability** -> **stable matching**.
- Questions:
 - Is stable matching possible?
 - If true, how to find it fast?



Possible – Proven with Examples

- Two men m , m' and two women w , w' .
 - m prefers w to w'
 - m' prefers w to w'
 - w prefers m to m'
 - w' prefers m to m'
- Do we have complete agreement?
 - The men agree on the order of women.
 - The women agree on the order of men.
 - Unique stable matching:** (m, w) and (m', w') .
 - Although m' prefers w , w didn't prefer m'
- Another perfect matching (m, w') and (m', w) .
 - Is this stable? **No (m and w prefer each other)**
- A more complicated example in the next slide.



Focus is
the pair
where they
want each
other
 (m, w)


$n = 2$, n^2 possibilities, $2^2=4$
 (m, w) , (m, w') , (m', w) , (m', w')
 6 possible combinations
 (m, w) , (m, w') – not possible
 (m, w) , (m', w) – not possible
 (m, w) , (m', w') - ok
 (m, w') , (m', w) – ok, but **unstable**
 (m, w') , (m', w') – not possible
 (m', w) , (m', w') – not possible

m and w wants each other
but now they are both
refused their choice and with
others

We have a situation where **2**
parties want each other but
yet they have someone else

Possible – Proven with Examples

- A more complicated example.



- m prefers w to w'
- m' prefers w' to w
- w prefers m' to m
- w' prefers m to m'
- Men's preferences meshes well $\rightarrow (m, w)$ and (m', w') perfect for men.
- Women's preferences meshes well $\rightarrow (m, w')$ and (m', w) perfect for women.
- Any problem?
 - Both are **stable matchings**, but they **clash**.
 - No happy ending for all men and women.
- Conclusion: ≥ 1 **stable matching possible**.

No situation where both parties want each other but yet get someone else

$(m, w), (m, w'), (m', w'), (m', w)$
 6 possible combinations
 $(m, w), (m, w')$ – not possible
 $(m, w), (m', w')$ – ok (fr men)
 $(m, w), (m', w)$ – not possible
 $(m, w'), (m', w')$ – not possible
 $(m, w'), (m', w)$ – ok (fr woman)
 $(m', w'), (m', w)$ – not possible

Design the Algorithm

- Our plan:
 - Show that a **stable matching exists**.
 - Propose an **efficient algorithm** accordingly.
- Some basic ideas – 3 Stages.
 - **Free**. In the beginning.
 - **Engagement**. Intermediate state. (**Still can change**)
 - **Married**. The algorithm or matching **terminates**. Engagements declared final.
- Why engage?
 - Bob likes Alice the most and proposes to her.
 - Shall Alice accept?
 - Accepts: what if Alice loves Jon the most and Jon proposes to her later? Regret?
 - If not, what if Bob is the best she can get? Lonely in the end?
 - **Engagement** is helpful to deal with the above situation.
- Procedure: all **free in the beginning**, **eventually get engaged**, and finally married.

Gale-Shapley Algorithm

Initially all $m \in M$ and $w \in W$ are free

While there is a man m who is free and hasn't proposed to every woman

 Choose such a man m

 Let w be the highest-ranked woman in m 's preference list to whom m has not yet proposed

 If w is free then

(m, w) become engaged

 Else w is currently engaged to m'

 If w prefers m' to m then

m remains free

 Else w prefers m to m'

(m, w) become engaged

m' becomes free

 Endif

Endif

Endwhile

Return the set S of engaged pairs

- Can return a **stable matching**?
- How to **prove** our judgement?

Woman w will become engaged to m if she prefers him to m' .




Figure 1.2 An intermediate state of the G-S algorithm when a free man m is proposing to a woman w .

Demo 1

Iteration	Man Proposes to	Woman	Engagement	Free Males	Reason
0				A, B, C	
1					
2					
3					
4					
5					
6					

Man	Preference List	Woman	Preference List
A	w_1, w_2, w_3	w_1	C, B, A
B	w_1, w_2, w_3	w_2	C, B, A
C	w_1, w_2, w_3	w_3	C, B, A

Demo 1 - Result

Male proposes, but woman get to select the most preferred man

Iteration	Man Proposes to	Woman	Engagement	Free Males	Reason
0				A, B, C	
1	A	w_1	(A, w_1)	B, C	1 st proposal to w_1
2	B	w_1	(B, w_1)	A, C	w_1 likes B more than A
3	A	w_2	$(B, w_1), (A, w_2)$	C	1 st proposal to w_2
4	C	w_1	$(C, w_1), (A, w_2)$	B	w_1 likes C more than B
5	B	w_2	$(C, w_1), (B, w_2)$	A	w_2 likes B more than A
6	A	w_3	$(C, w_1), (B, w_2), (A, w_3)$	END	1 st proposal to w_3

Man	Preference List	Woman	Preference List
A	w_1, w_2, w_3	w_1	C, B, A
B	w_1, w_2, w_3	w_2	C, B, A
C	w_1, w_2, w_3	w_3	C, B, A

Demo 2 (Let Women propose for fairness)

Iteration	Wm Proposes to	Man	Engagement	Free Wms	Reason
0				A, B, C	
1					
2					
3					
4					
5					
6					

Wms	Preference List	Man	Preference List
A	m_1, m_2, m_3	m_1	A, B, C
B	m_1, m_2, m_3	m_2	A, B, C
C	m_1, m_2, m_3	m_3	A, B, C

Demo 2 - Result

Woman proposes, but man get to select the most preferred woman

Iteration	Wm Proposes to	Man	Engagement	Free Wms	Reason
0				A, B, C	
1	A	m_1	(A, m_1)	B, C	1 st proposal to m_1
2	B	m_1	(A, m_1)	B, C	m_1 (with A) rejects B
3	B	m_2	$(A, m_1), (B, m_2)$	C	1 st proposal to m_2 (SUC)
4	C	m_1	$(A, m_1), (B, m_2)$	C	m_1 (with A) rejects C
5	C	m_2	$(A, m_1), (B, m_2)$	C	m_2 (with B) rejects C
6	C	m_3	$(A, m_1), (B, m_2), (C, m_3)$	END	1 st proposal to m_3 (SUC)

Wms	Preference List	Man	Preference List
A	m_1, m_2, m_3	m_1	A, B, C
B	m_1, m_2, m_3	m_2	A, B, C
C	m_1, m_2, m_3	m_3	A, B, C

Analyze the Algorithm

- (Assume men proposing)
- **Fact 1.1.** Woman w remains engaged from the point at which she receives her first proposal; and the sequence of partners to which she is engaged gets better and better.
 - From the woman's perspective.
 - Only changes her mind if a more preferred man proposes.
- **Fact 1.2.** The seq. of women to whom man m proposes gets worse and worse.
 - From the man's perspective.
 - Only proposes again if he is free or his engaged more preferred woman left him.
- **Fact 1.3.** The algorithm terminates after at most $O(n^2)$ iterations.
 - Proof trick: focus on the terminate condition of the while loop.
 - One man proposes to one different woman (he never proposed to) in each iteration.
 - How many men? How many woman? $n \times n$

$n = 2$, n^2 possibilities, $2^2=4$
 $(m,w), (m,w'), (m',w), (m',w')$

Analyze the Algorithm

- **Fact 1.4.** If man m is free at some point in the execution of the algorithm, then there is a woman to whom he has not yet proposed.
 - Proof idea: Contradiction analysis. Suppose the man has proposed to all, but now he is free. The only possibility is that all his proposed women are engaged with other men.
 - We have n women, engaged to n men, which does not include the man m . Now, we have $n+1$ men. But it cannot be true.
- **Fact 1.5.** The set S returned at termination is a perfect matching.
 - Proof idea: The algorithm terminates only if there is no free man. So, all men and women are in the matching.

Analyze the Algorithm

- **Fact 1.6.** The algorithm **returns a stable matching**.
 - Proof idea: at least know it is perfect, so prove there is **no instability**.
 - **Contradiction analysis** again. Suppose we really have one:
 - (m, w) and (m', w') that:
 - Man m prefers w' more && woman w' prefers m more $\rightarrow (m, w')$?
 - If man m is together with w , he must have **proposed to w' before**, as he like w' more.
 - So w' **rejected m for m'** ? No way, as w' like m more than m' . (Fact 1.1)

Wms	Preference List	Man	Preference List
w		m	w', w
w'	m, m'	m'	

Computational Tractability



Computational Tractability

- **Resources** in computer systems?
 - **Time, space, communication.**
[<https://ieeexplore.ieee.org/document/9063714>]
 - Which one is the most important?
 - We **focus on time**, and **space** is also considered.
 - Why do we consider space?
- Obj: **quantify computational tractability**. How it changes?
 - Proposing an algorithm is easy.
 - An **efficient** (and correct) algorithm will be challenging.

<https://analyticsindiamag.com/wp-content/uploads/2019/10/computer-types-10.jpg>

https://i.pcmag.com/imagery/lineups/02SYcRWQLDk6G1bnereOVv-2.fit_lim.v_1585232946.jpg

https://miro.medium.com/max/600/0*AgIPFgJrm9GQfzKJ.jpg

Computational Tractability

- Informal **efficiency** def.: **runs quickly on real input**.
 - Problem 1: run the algorithm in what **environment**?
 - Problem 2: how **fast** is fast?
 - Special case: a poor-designed algorithm runs **super fast** with **tiny data input**
 - Special case: a well-designed algorithm runs **super slow** with **huge data**
 - Problem 3: what is **real input**? Crash with large input?
- We want: **platform-independent** and **instance-independent**.



<https://analyticsindiamag.com/wp-content/uploads/2019/10/computer-types-10.jpg>

https://i.pcmag.com/imagery/lineups/02SYcRWcQLDk6G1bnereOVv-2.fit_lim.v_1585232946.jpg

https://miro.medium.com/max/600/0*AgIPFgJrm9GQfzKJ.jpg

Worst-Case Running Time

- Quantify the input size: N
 - n men and n women (eg. $n=3$)
 - $2n$ preference lists, each list of length n (eg. 2×3 plist $\times 3$)
 - $N = 2n \cdot n = 2n^2$
- Task: analyze the running time as a function of N .
- Worst-case running time: the largest possible running time with input N .
- Benchmark: brute-force algorithm (brainless design).

Man	Preference List	Woman	Preference List
A	w_1, w_2, w_3	w_1	C, B, A
B	w_1, w_2, w_3	w_2	C, B, A
C	w_1, w_2, w_3	w_3	C, B, A

Redefine Efficiency

- Informal efficiency def.: achieves better worst-case performance, at an analytical level, than brute-force search.
- **Polynomial time**: quantify “reasonable” running time.
 - **Search space**: often grows exponentially.
 - **Reasonable** may be: from n to $n + 1$, running time increases by a constant.
- Informal efficiency def.: an alg. is efficient if runs in polynomial time.
 - Extreme cases exists like the worst-case concept, i.e., $n^{10000000}$?
 - Again, in reality, the polynomial concept works well.

Running Times

n is the
number
of data
points


	n	$n \log_2 n$	n^2	n^3	1.5^n	2^n	$n!$
$n = 10$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	4 sec
$n = 30$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	< 1 sec	18 min	10^{25} years
$n = 50$	< 1 sec	< 1 sec	< 1 sec	< 1 sec	11 min	36 years	very long
$n = 100$	< 1 sec	< 1 sec	< 1 sec	1 sec	12,892 years	10^{17} years	very long
$n = 1,000$	< 1 sec	< 1 sec	1 sec	18 min	very long	very long	very long
$n = 10,000$	< 1 sec	< 1 sec	2 min	12 days	very long	very long	very long
$n = 100,000$	< 1 sec	2 sec	3 hours	32 years	very long	very long	very long
$n = 1,000,000$	1 sec	20 sec	12 days	31,710 years	very long	very long	very long

- Assumptions:
 - 1m instructions per second
 - Very long \rightarrow runs over 10^{25} years.
- Intelligence does not necessarily produce polynomial algorithms.

1,000,000 instructions per second (1mins = 60secs)
 $n^3 = 1,000^3 = 1,000,000,000$
 $1,000,000,000 / 1,000,000 = 1,000\text{secs} \times 1/60 = 17\text{mins}$

Asymptotic Order of Growth – Upper Bound


- Target: **worst-case** time, given input n , grows at a rate **at most proportional to $f(n)$** .
 - We also call $f(n)$ a **bound**.
- **Too specific expressions unnecessary**, i.e., runs in $1.62n^2 + n + 4$ steps.
 - Constants vary for different architectures
 - Determined by the highest order (e.g. n^2)
- How do you plan to quantify the above terms concisely?
 - We **ignore the constants** and We **ignore the lower order items**.
- **Asymptotic Upper Bounds $O(\cdot)$** :
 - The **worst-case running time $T(n)$** of a certain algorithm on an input of size n .
 - Expresses only an **upper bound**.
 - $f(n) = pn^2 + qn + r$ for constants p , q , and $r \rightarrow O(n^2)$.




Asymptotic Order of Growth – More bounds

- How to justify our proven bound $O(\cdot)$ is good?
 - i.e., we are almost always true if we say it is $O(n^{100000n})$.
- **Asymptotic Lower Bounds : $\Omega(n)$.**
 - The func. $T(n)$ is **at least** a constant “c” multiple of $f(n) = c \cdot g(n)$.
 - $T(n) = pn^2 + qn + r$ for constants p, q , and $r \rightarrow \Omega(n^2)$.
- **Asymptotically Tight Bounds $\Theta(n)$:** $f(n) = \Omega(g(n)) = O(g(n))$ (optimal bound)
 - i.e., $f(n) = pn^2 + qn + r = \Omega(n^2) = O(n^2) = \Theta(n^2)$.
- Note: consider $n \rightarrow \infty$.
 - Only matters when n is large.
 - $f(n) = \Theta(g(n))$ if $\lim_{n \rightarrow \infty} \frac{f(n)}{g(n)} = c$.
 - c is a positive constant.


$c \cdot g(n)$ is the estimated bound



(a) Big-O Notation




(b) Big-Omega Notation




(c) Big-Theta Notation

- $f(n) = 2n + 3$
- $f(n) = O(g(n))$ where $g(n) = n$
- $f(n) \leq c_1 \cdot g(n) = c_1 \cdot n$ where $n > n_0$, $c_1 > 0$, $n_0 \geq 1$
- $2n+3 \leq c_1 \cdot n$ where the RHS, big-O is larger than actual function
- If $n_0=1$ and $c_1 = 5$, then $5 \leq 5$: ok
- If $n_0=2$ and $c_1 = 5$, then $7 \leq 10$: ok
- Conclude that the big-oh of the function $2n+3$ is $O(n)$ where $c_1=5$, $n_0 \geq 1$. This is the best worst case of this algorithm
- $f(n) = \Omega(g(n))$ where $g(n) = n$
- $f(n) \geq c_2 \cdot g(n) = c_2 \cdot n$ where $n > n_0$, $c_2 > 0$, $n_0 \geq 1$
- $2n+3 \geq c_2 \cdot n$ where RHS, big-omega is smaller than the actual function
- If $n_0=1$ and $c_2 = 1$, then $5 \geq 1$: ok
- If $n_0=2$ and $c_2 = 1$, then $7 \geq 2$: ok
- Conclude that the big-omega of the function $2n+3$ is $\Omega(n)$ where $c_2=1$, $n_0 \geq 1$. This is the best case of this algorithm (the algo will never be better than this)



- $f(n) = 2n + 3$
- $f(n) = \theta(g(n))$ where $g(n) = n$
- $c_2 \cdot n = c_2 \cdot g(n) \leq f(n) \leq c_1 \cdot g(n) = c_1 \cdot n$ where $n > n_0$, $c_1, c_2 > 0$, $n_0 \geq 1$
- $2n+3 \leq c_1 \cdot n$ where the RHS is the worst case
- Conclude that the worst case of the function $2n+3$ is $\theta(n)$ where $c_1=5$, $n_0 \geq 1$. This is the best worst case of this algorithm
- $c_2 \cdot n \leq 2n+3$ where LHS, is the best case
- Conclude that the best case of the function $2n+3$ is $\theta(n)$ where $c_2=1$, $n_0 \geq 1$. This is the best case of this algorithm (the algo will never be better than this)
- **theta is the average case of the algorithm**




Common Times - $O(n)$

- Long-time goal: recognize **common styles** of analyzing the running time.
- **Linear time**: $O(n)$, time scales by a constant factor times the input size n .
- **Find the maximum**: a_1, a_2, \dots, a_n .
 - Each needs a comparison, in constant time.
 - $O(n) * O(1) = O(n)$.
- **Merge sorted lists**. The same $O(n)$, but more complex.
 - List a_1, a_2, \dots, a_n and list b_1, b_2, \dots, b_n , both in ascending order.
 - Objective: c_1, c_2, \dots, c_{2n} .
 - i.e., 2, 3, 11, 19 and 4, 9, 16, 25 \rightarrow 2, 3, 4, 9, 11, 16, 19, 25.
 - A non-intelligent solution: concatenate the lists \rightarrow sort all the items.
 - Record pa , pb , and pc , where $pa = pb = pc = 1$ in the beginning.
 - While $pa \leq n$ and $pb \leq n$:
 - If $a_{pa} > b_{pb} \rightarrow$ smaller b_{pb} to $c_{pc} \rightarrow pb++$ and $pc++$.
 - If $a_{pa} \leq b_{pb} \rightarrow$ smaller a_{pa} to $c_{pc} \rightarrow pa++$ and $pc++$.
 - Whatever left, append to c .

```

max = a1
For i = 2 to n
    If ai > max then
        set max = ai
    Endif
Endfor
  
```

Append the smaller of a_i and b_j to the output.



Common Times - $O(n \log n)$, $O(n^2)$, $O(n^3)$, $O(n^k)$



- **$O(n \log n)$** : the time of any alg. **splits input into equal-sized pieces**, **solves each recursively**, and **combines** the two **solutions** in linear time.
 - Popular example: **sorting**.
 - Analyze the basic idea: sorting is expensive \rightarrow sort less \rightarrow combine.
 - i.e., $100^2 = 10,000 \rightarrow 2 \times 50^2 = 5,000 \rightarrow 5,000 + 100 = 5,100$, around half time.
- **$O(n^2)$** : i.e., **n points in a plane**, what the **nearest points**?
 - Brute-force sol.: for each pt, compute the **dist. between it and another pt**.
 - $O(n)$ points, each $O(n)$ distances, each distance $O(1) \rightarrow O(n^2)$.
- **$O(n^3)$** : nested loops, i.e., **3-variable solver**: $2x + 5y + 9z = 100$.
 - For $x \rightarrow$ for $y \rightarrow$ for $z \rightarrow$ check the sum.
- **$O(n^k)$** : Any idea on getting $O(n^k)$?
 - Independent sets in a graph with k nodes.
 - Brute-force, **choose k from n** : $\binom{n}{k} = \frac{n(n-1)\cdots(n-k+1)}{1 \cdot 2 \cdots k} \leq \frac{n^k}{k!} \leq n^k$.

Beyond Polynomial Time

- $O(2^n)$ vs. $O(n!)$. Which one is faster?
 - $O(2^n)$: the number of the subsets of n nodes.
 - $O(n!)$: the number of perfect matchings; traveling salesman problem.
- **Sublinear time**: better than linear, i.e., $O(\log n)$.
 - Binary search.
 - **Query** the inputs instead of traversing them.
 - Half \rightarrow half \rightarrow half \rightarrow half ...